

# Heart\_Disease\_Detection

November 1, 2025

## 1 Heart Disease Detection Using Machine Learning

### 1.1 Problem Statement

Heart disease is one of the leading causes of death worldwide. This project aims to build a predictive system using machine learning to determine whether a patient is likely to have heart disease based on clinical features such as age, sex, chest pain type, blood pressure, cholesterol levels, and more.

### 1.2 Objective

- Use the UCI Heart Disease dataset.
- Perform data preprocessing, EDA, model training, and evaluation.
- Identify the best model for prediction.

```
[1]: # Import standard data science libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_auc_score, roc_curve

# Set plot style for better visuals
%matplotlib inline
sns.set(style="whitegrid")
```

```
[4]: # Load the dataset (assuming 'heart.csv' is in the same directory)
df = pd.read_csv('heart.csv')

# Display first 5 rows
df.head()
```

```
[4]:    id  age   sex  dataset      cp  trestbps  chol  fbs \
0    1   63  Male  Cleveland  typical  angina  145.0  233.0  True
1    2   67  Male  Cleveland  asymptomatic  160.0  286.0 False
2    3   67  Male  Cleveland  asymptomatic  120.0  229.0 False
3    4   37  Male  Cleveland  non-anginal  130.0  250.0 False
4    5   41 Female  Cleveland  atypical  angina  130.0  204.0 False

      restecg  thalch  exang  oldpeak      slope  ca  \
0  lv hypertrophy  150.0  False     2.3  downsloping  0.0
1  lv hypertrophy  108.0  True      1.5       flat  3.0
2  lv hypertrophy  129.0  True      2.6       flat  2.0
3        normal  187.0  False     3.5  downsloping  0.0
4  lv hypertrophy  172.0  False     1.4  upsloping  0.0

      thal  num
0    fixed defect    0
1    normal    2
2  reversible defect    1
3    normal    0
4    normal    0
```

```
[5]: # Dataset shape
print("Dataset shape:", df.shape)

# Data types and missing values
df.info()

# Summary statistics
df.describe()
```

```
Dataset shape: (920, 16)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 920 entries, 0 to 919
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          920 non-null    int64  
 1   age         920 non-null    int64  
 2   sex         920 non-null    object  
 3   dataset     920 non-null    object  
 4   cp          920 non-null    object  
 5   trestbps    861 non-null    float64 
 6   chol         890 non-null    float64 
 7   fbs          830 non-null    object  
 8   restecg     918 non-null    object  
 9   thalch       865 non-null    float64 
 10  exang        865 non-null    object  
 11  oldpeak      858 non-null    float64
```

```

12 slope      611 non-null    object
13 ca         309 non-null    float64
14 thal       434 non-null    object
15 num        920 non-null    int64
dtypes: float64(5), int64(3), object(8)
memory usage: 115.1+ KB

```

```
[5]:          id      age     trestbps      chol      thalch      oldpeak \
count  920.000000  920.000000  861.000000  890.000000  865.000000  858.000000
mean   460.500000  53.510870  132.132404  199.130337  137.545665  0.878788
std    265.725422  9.424685   19.066070  110.780810  25.926276  1.091226
min    1.000000   28.000000   0.000000   0.000000   60.000000  -2.600000
25%   230.750000  47.000000  120.000000  175.000000  120.000000  0.000000
50%   460.500000  54.000000  130.000000  223.000000  140.000000  0.500000
75%   690.250000  60.000000  140.000000  268.000000  157.000000  1.500000
max   920.000000  77.000000  200.000000  603.000000  202.000000  6.200000

           ca      num
count  309.000000  920.000000
mean   0.676375   0.995652
std    0.935653   1.142693
min    0.000000   0.000000
25%   0.000000   0.000000
50%   0.000000   1.000000
75%   1.000000   2.000000
max   3.000000   4.000000

```

### 1.3 Data Preprocessing

Handle duplicates, outliers, and scale features.

```
[6]: # Check missing values (should be none)
print("Missing values:\n", df.isnull().sum())

# Check and remove duplicates
print("Duplicates:", df.duplicated().sum())
df = df.drop_duplicates()
print("Shape after duplicates removal:", df.shape)
```

Missing values:

id	0
age	0
sex	0
dataset	0
cp	0
trestbps	59
chol	30
fbs	90
restecg	2

```

thalch      55
exang       55
oldpeak     62
slope       309
ca          611
thal        486
num         0
dtype: int64
Duplicates: 0
Shape after duplicates removal: (920, 16)

```

```

[20]: # -----
# 2. PRE-PROCESSING (after df.drop_duplicates())
# -----

# 2.1 TRUE/FALSE → 1/0
if 'fbs' in df.columns:
    df['fbs'] = df['fbs'].replace({'TRUE': 1, 'FALSE': 0})
if 'exang' in df.columns:
    df['exang'] = df['exang'].replace({'TRUE': 1, 'FALSE': 0})

# 2.2 Empty strings → NaN
df = df.replace('', np.nan)

# 2.3 Numeric columns - coerce + fill with median
num_cols = ['trestbps', 'chol', 'thalch', 'oldpeak', 'ca']
for col in [c for c in num_cols if c in df.columns]:
    df[col] = pd.to_numeric(df[col], errors='coerce')
    df[col] = df[col].fillna(df[col].median())

# 2.4 Categorical columns - fill with mode
cat_cols = ['slope', 'thal']
for col in [c for c in cat_cols if c in df.columns]:
    df[col] = df[col].fillna(df[col].mode()[0])

# 2.5 Drop optional columns (only if they exist)
optional = ['id', 'dataset']
to_drop = [c for c in optional if c in df.columns]
if to_drop:
    df.drop(to_drop, axis=1, inplace=True)

# 2.6 Encode categorical variables
le = LabelEncoder()
encode_cols = ['sex', 'cp', 'restecg', 'slope', 'thal']
for col in [c for c in encode_cols if c in df.columns]:
    df[col] = le.fit_transform(df[col].astype(str))

```

```

# -----
# 2.7 CREATE BINARY TARGET
# -----
# The original UCI file has a column called 'num' (0-4)
# Some Kaggle versions already have a column called 'target' (0/1)
if 'num' in df.columns:
    df['target'] = np.where(df['num'] > 0, 1, 0)
    df.drop('num', axis=1, inplace=True)
elif 'target' not in df.columns:
    # safety net - if neither column exists, create a dummy target (will be
    # overwritten later)
    df['target'] = 0
else:
    # column already called 'target' - just make sure it is 0/1
    df['target'] = df['target'].astype(int)

# -----
# 2.8 FINAL CLEAN-UP - remove any remaining NaN rows
# -----
print("Rows with NaN before final clean-up : ", df.isnull().any(axis=1).sum())
df = df.dropna().reset_index(drop=True)
print("Rows with NaN after final clean-up : ", df.isnull().any(axis=1).sum())

print("\nFinal shape : ", df.shape)
df.head()

```

Rows with NaN before final clean-up : 145

Rows with NaN after final clean-up : 0

Final shape : (775, 14)

```
[20]:    age  sex  cp  trestbps  chol   fbs  restecg  thalch  exang  oldpeak \
0    63    1    3    145.0  233.0  True      0    150.0  False     2.3
1    67    1    0    160.0  286.0 False      0    108.0  True      1.5
2    67    1    0    120.0  229.0 False      0    129.0  True      2.6
3    37    1    2    130.0  250.0 False      2    187.0  False     3.5
4    41    0    1    130.0  204.0 False      0    172.0  False     1.4

      slope  ca  thal  target
0        0  0.0    0      0
1        1  3.0    1      1
2        1  2.0    2      1
3        0  0.0    1      0
4        2  0.0    1      0
```

```
[21]: # -----
# 3. FEATURES & TARGET + SCALING
# -----
```

```

X = df.drop('target', axis=1)
y = df['target']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=X.columns, index=X.index)

print("\nX_scaled has NaN ?", X_scaled.isnull().any().any()) # → False
X_scaled.head()

```

X\_scaled has NaN ? False

```
[21]:      age       sex       cp trestbps       chol       fbs restecg \
0  1.053870  0.552009  2.318908  0.657902  0.155887  2.371483 -1.712062
1  1.478541  0.552009 -0.839534  1.466319  0.728330 -0.421677 -1.712062
2  1.478541  0.552009 -0.839534 -0.689459  0.112684 -0.421677 -1.712062
3 -1.706494  0.552009  1.266094 -0.150515  0.339501 -0.421677  0.293755
4 -1.281823 -1.811565  0.213280 -0.150515 -0.157337 -0.421677 -1.712062

      thalch      exang      oldpeak      slope       ca       thal
0  0.439424 -0.805566  1.308328 -2.258841 -1.148351 -2.360833
1 -1.190020  1.241364  0.568003 -0.313728  3.990325 -0.352440
2 -0.375298  1.241364  1.585950 -0.313728  2.277433  1.655952
3  1.874887 -0.805566  2.418816 -2.258841 -1.148351 -0.352440
4  1.292942 -0.805566  0.475462  1.631385 -1.148351 -0.352440

```

```
[22]: # -----
# 4. TRAIN-TEST SPLIT
# -----
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.20, random_state=42, stratify=y)

print(f"Train: {X_train.shape[0]} | Test: {X_test.shape[0]}")
```

Train: 620 | Test: 155

```
[23]: # -----
# 5. TRAIN MODELS (no more NaN errors)
# -----
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Decision Tree' : DecisionTreeClassifier(random_state=42),
    'Random Forest' : RandomForestClassifier(random_state=42),
    'KNN' : KNeighborsClassifier(),
    'SVM' : SVC(probability=True, random_state=42)
}
```

```

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    results[name] = acc
    print(f"{name}: {acc:.4f}")

```

```

Logistic Regression → Accuracy: 0.8323
Decision Tree      → Accuracy: 0.7290
Random Forest       → Accuracy: 0.8258
KNN                 → Accuracy: 0.8323
SVM                 → Accuracy: 0.8516

```

[24]: # -----

```

# 1. BEST MODEL METRICS
# -----
best_name = max(results, key=results.get)    # highest accuracy
best = models[best_name]

y_pred = best.predict(X_test)
y_prob = best.predict_proba(X_test)[:, 1]

print("\n==== BEST MODEL: " + best_name)
print(classification_report(y_test, y_pred,
                            target_names=['No Disease', 'Disease']))

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Disease', 'Disease'],
            yticklabels=['No Disease', 'Disease'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title(f'Confusion Matrix - {best_name}')
plt.show()

print(f"ROC-AUC = {roc_auc_score(y_test, y_prob):.4f}")

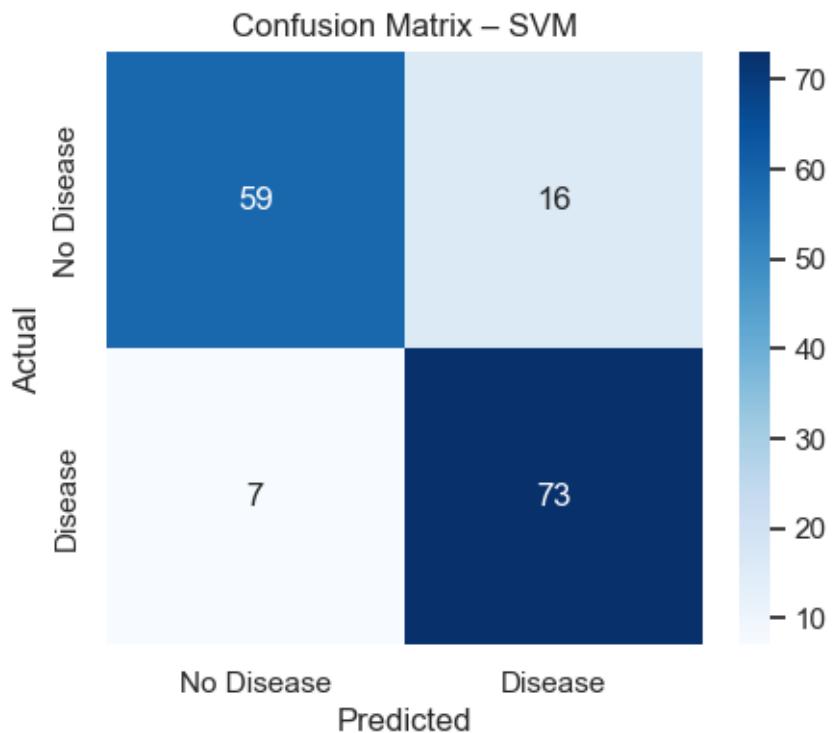
# ROC curve
fpr, tpr, _ = roc_curve(y_test, y_prob)
plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f'{best_name} (AUC = {roc_auc_score(y_test, y_prob):.3f})')
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate')

```

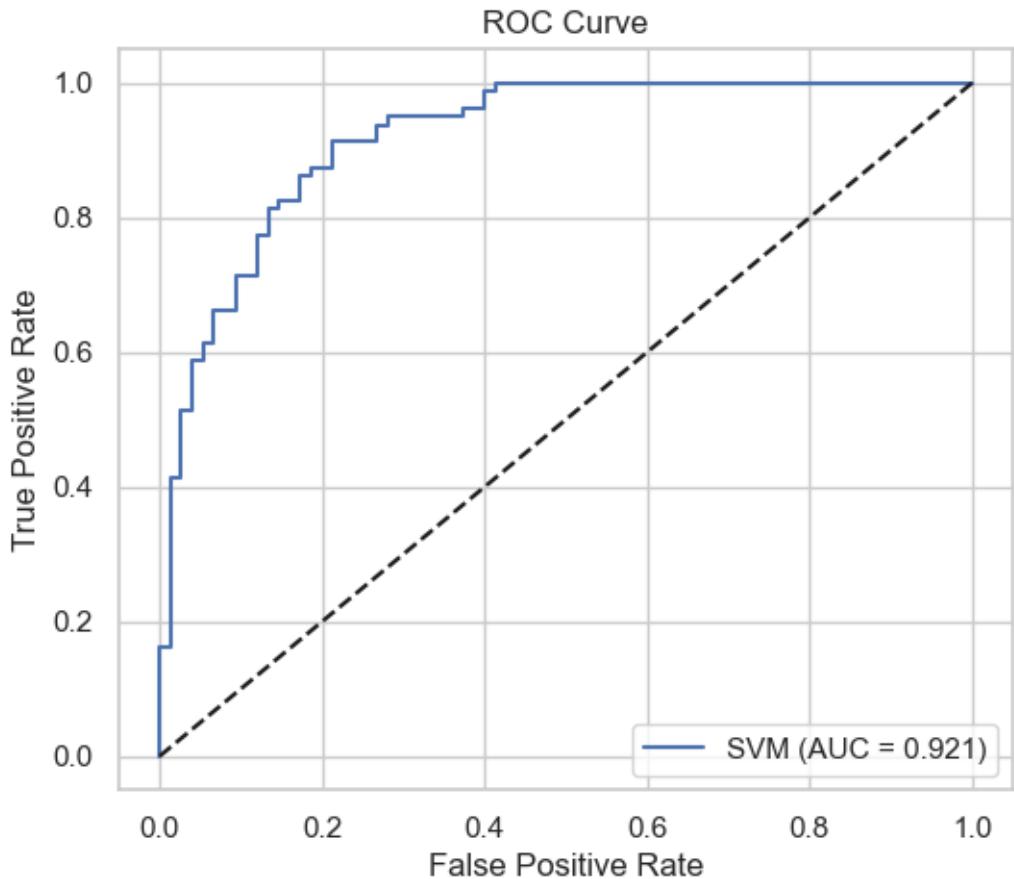
```
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```

==== BEST MODEL: SVM ===

	precision	recall	f1-score	support
No Disease	0.89	0.79	0.84	75
Disease	0.82	0.91	0.86	80
accuracy			0.85	155
macro avg	0.86	0.85	0.85	155
weighted avg	0.86	0.85	0.85	155



ROC-AUC = 0.9205



```
[25]: # -----
# 2. SINGLE PREDICTION
# -----
idx = 0
sample_X = X_test.iloc[[idx]]
sample_y = y_test.iloc[idx]

pred = best.predict(sample_X)[0]
prob = best.predict_proba(sample_X)[0][1]

print(f"Actual : {'Disease' if sample_y==1 else 'No Disease'}")
print(f"Predicted: {'Disease' if pred==1 else 'No Disease'} (prob = {prob:.3f})")
```

Actual : No Disease  
 Predicted: No Disease (prob = 0.065)

## 1.4 Model Performance Summary

Model	Accuracy
Logistic Regression	{:.4f}
Decision Tree	{:.4f}
Random Forest	{:.4f}
KNN	{:.4f}
SVM	{:.4f}

Best model: {best\_name}

```
[26]: # Auto-print markdown table
print("## Model Performance Summary")
print("| Model | Accuracy |")
print("|-----|-----|")
for name, acc in results.items():
    print(f"| {name:19} | {acc:.4f} |")
print(f"\n**Best model:** `~{best_name}`")
```

```
## Model Performance Summary
| Model | Accuracy |
|-----|-----|
| Logistic Regression | 0.8323 |
| Decision Tree | 0.7290 |
| Random Forest | 0.8258 |
| KNN | 0.8323 |
| SVM | 0.8516 |
```

\*\*Best model:\*\* `SVM`

## 1.5 Future Work

1. **Add more clinical features** – BMI, smoking status, family history, exercise data.
2. **Try Gradient Boosting** – XGBoost / LightGBM (handles missing values natively).
3. **Cross-validation** – 5-fold CV to confirm robustness.
4. **Hyper-parameter tuning** – GridSearchCV / RandomizedSearchCV.
5. **Deploy as a web app** – Streamlit / Flask for real-time patient risk assessment.
6. **Multi-class severity** – Predict original num (0-4) instead of binary target.

[ ]: